

Assignment 5

Deadline: December 11, 23:59

Assignment instructions:

The assignment is to be done in **pairs**. You'll need to join a **group on Canvas** to be able to submit a solution for grading.

Individual or late submissions will not be graded.

Your submission will consist of the following components:

1. A **PDF-report**. It's highly recommended for you to write your report in Latex, using the provided template on Canvas. Hand-written submissions will not be graded unless they're *very very* easy to read. You run the risk of losing unnecessary points if you submit anything hand-written.
2. Your **code**. This can be included in **either** an appendix, or a standalone file-upload. Please do not upload a zip-file containing multiple files. Gather all your code in one .py or .ipynb in case you don't wish to add it to the appendix.

Figures asked for in the programming exercises should be well presented, meaning that they have: high image quality, a title and a caption, labeled axes and a legend for the curves if needed.

In case you run into problems, the consultation hours are on Wednesdays 13:15-15:00. If possible, send an email with your questions to book a time with the responsible TA. The responsible TA for the assignment is the one who uploaded it to Canvas.

There's more information regarding the assignments in canvas, both in the standalone module, and in the slides of tutorial session 1.

Best of luck!

Exercise 1: Gibbs Sampling (50%)

1. Derive the Conditional Distributions

For the Gaussian model:

$$y_i \sim \mathcal{N}(\mu, \tau^{-1}),$$

where $\mu \sim \mathcal{N}(0, \omega^{-1})$ and $\tau \sim \text{Gamma}(\alpha, \beta)$.

Q1: Prove the following:

$$\mu|y, \tau \sim \mathcal{N}\left(\frac{\tau}{n\tau + \omega} \sum y_i, \frac{1}{n\tau + \omega}\right)$$

$$\tau|y, \mu \sim \text{Gamma}\left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum (y_i - \mu)^2\right).$$

For the Poisson model:

$$y_i \sim \text{Poisson}(\mu),$$

where $\mu \sim \text{Gamma}(2, \beta)$ and $\beta \sim \text{Exponential}(1)$.

Q2: Prove the following:

$$\mu|y, \beta \sim \text{Gamma}\left(2 + \sum_i y_i, n + \beta\right)$$

$$\beta|y, \mu \sim \text{Gamma}(3, 1 + \mu).$$

2. Implement Gibbs Sampler

Q3: Gaussian Model

- Using the provided Python code (`E1.ipynb`) as a base, implement Gibbs sampling for the Gaussian model.
- Simulate $n = 100$ observations from $y_i \sim \mathcal{N}(5, 2^{-1})$.
- Plot traceplots and histograms for μ and τ^{-1} (variance). Compare with true values.

Q4: Poisson Model

- Extend the Gibbs sampler to the Poisson model with $n = 100$ observations from $y_i \sim \text{Poisson}(5)$.
- Plot traceplots and histograms for μ and β . Compare with theoretical expectations.

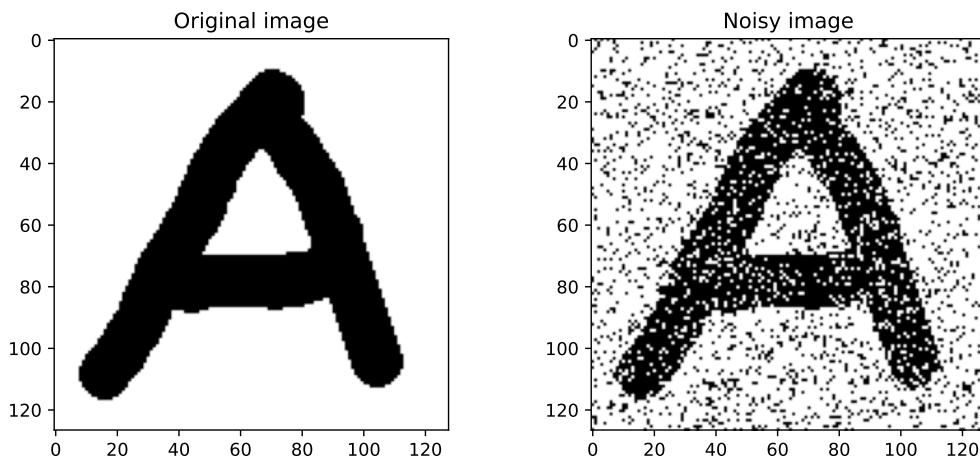


Figure 1: Original vs noisy image

Exercise 2: Minimization of Energy (50%)

One application of Markov Random Fields (MRFs) is to denoise noisy images. In this task, we aim to minimize the following energy function using **Gibbs Sampling** and **Mean Field Approximation**:

$$E(x, y) = h \sum_i x_i - \beta \sum_{i,j} x_i x_j - \eta \sum_i x_i y_i,$$

where:

- $x_i \in \{-1, +1\}$: Binary state of pixel i in the denoised image,
- $y_i \in \{-1, +1\}$: Binary state of pixel i in the noisy observed image,
- h , β , and η : Scalar parameters controlling the trade-offs between the terms.

You can download the clear image for this exercise from the assignment page. The image is stored in a CSV file format. You can read it into Python as follows:

```
import pandas as pd
img = pd.read_csv('letterA.csv').to_numpy()
```

Algorithms

Gibbs Sampling

1. **Initialization:** Set $x_i = y_i$ for all i .
2. For each pixel $j = 1, \dots, N$:

- Compute the conditional probability:

$$P(x_j = +1 \mid x_{\setminus j}, y) \propto \exp(-E(x, y)),$$

where $x_{\setminus j}$ denotes all pixel states except x_j .

- Sample x_j from the distribution:

$$P(x_j \mid x_{\setminus j}, y).$$

3. Repeat for multiple iterations or until convergence.

Mean Field Approximation

Now, we try to apply mean-field approximation to denoise the image. The posterior for the Ising model has the following form:

$$P(x \mid y) = \frac{1}{Z} \exp(-E(x)),$$

where:

$$E(x) = E_0(x) - \sum_i L_i(x_i), \quad E_0(x) = -\beta \sum_{i,j} x_i x_j, \quad L_i(x_i) = \eta \sum_i x_i y_i.$$

In mean-field approximation, we try to approximate this posterior by a fully factored approximation:

$$q(x) = \prod_i q(x_i; \mu_i),$$

where μ_i is the mean value of node i . To derive the update for the variational parameter μ_i , we first write out:

$$\log \tilde{p}(x) = -E(x),$$

dropping terms that do not involve x_i :

$$\log \tilde{p}(x) = \beta x_i \sum_j x_j + L_i(x_i) + \text{const.}$$

This only depends on the states of the neighboring nodes. Now we take expectations of this with respect to $\prod_{j \neq i} q(x_j)$ to get:

$$q(x_i) \propto \exp \left(x_i \sum_j \beta \mu_j + L_i(x_i) \right).$$

Thus, we replace the states of the neighbors by their average values. Let:

$$m_i = \sum_j \beta \mu_j,$$

be the mean field influence on node i . Also, let $L_i^+ = L_i(1)$ and $L_i^- = L_i(-1)$. The approximate marginal posterior is given by:

$$q_i(x_i = 1) = \frac{\exp(m_i + L_i^+)}{\exp(m_i + L_i^+) + \exp(-m_i + L_i^-)} = \sigma(2a_i),$$

where:

$$a_i = m_i + 0.5(L_i^+ + L_i^-).$$

Similarly, we have $q_i(x_i = -1) = \sigma(-2a_i)$. From this, we can compute the new mean for site i :

$$\mu_i = \mathbb{E}_{q_i}[x_i] = \tanh(a_i).$$

Hence, the update equation becomes:

$$\mu_i = \tanh \left(\sum_j \beta \mu_j + 0.5(L_i^+ - L_i^-) \right).$$

It is usually better to use damped updates of the form:

$$\mu_i^t = (1 - \lambda)\mu_i^{t-1} + \lambda \tanh \left(\sum_j \beta \mu_j^{t-1} + 0.5(L_i^+ - L_i^-) \right), \quad 0 < \lambda < 1.$$

Tasks

Q1 Implement Gibbs Sampling for Image Denoising:

- Simulate noisy images with varying noise levels (e.g., Gaussian noise or salt-and-pepper noise).
- Implement Gibbs Sampling to minimize the energy function.
- Apply the algorithm to the noisy image and produce a denoised image.
- Plot the noisy image and the denoised images after iterations from the following list: [0, 1, 5, 15, 50].

Q2 Implement Mean Field Approximation for Image Denoising:

- Implement the mean-field approximation using the update equations provided.
- Apply the algorithm to the noisy images and produce a denoised image.
- Plot the noisy and denoised image after iterations from the following list: [0, 1, 5, 15, 50].

Q3 Noise and Performance Analysis:

- Compute and plot the **Normalized Mean Squared Error** (NMSE) between the denoised image (D) and the original ground truth image (G) for each noise level:

$$\text{NMSE} = \frac{\|D - G\|^2}{\|G\|^2}.$$

- Compare the performance of Gibbs Sampling and Mean Field Approximation as the noise level increases.

Q4 Comparison of Results:

- Evaluate the trade-offs between the methods, considering:
 - Convergence speed,
 - Final NMSE values,
 - Visual quality of the denoised image.

Q5 Parameter Sensitivity Analysis:

- Vary the parameters β , and η in the energy function and evaluate their impact on denoising performance for both algorithms.
- For each parameter, plot:
 - NMSE vs β ,
 - NMSE vs η .
- Discuss and justify how each parameter influences the trade-offs between pixel similarity, smoothness, and adherence to the noisy image.